

Processes & Threads

Processes

- **Process:** Program in execution.
- **Multiprogramming:** Concurrent processes on one core.
- **Multiprocessing:** One process on multiple cores.
- **Address Space:** Code (Text), Static Data, Heap, Stack.
- **PCB (Process Control Block):** `task_struct`, pointed to by `current`.
- **PCB Fields:** PID, TGID, state, registers, `mm_struct`, family pointers, exit code.
- **Process States:**
 - **Running:** On CPU.
 - **Ready:** Waiting for CPU.
 - **Waiting (Sleeping):** For event. `TASK_INTERRUPTIBLE`, `TASK_UNINTERRUPTIBLE`.
 - **Stopped:** Suspended.
 - **Zombie:** Terminated, PCB exists for `wait()`.
- **Process Family:** Idle (PID 0) → `init` (PID 1). Real Parent, Parent, Orphans (adopted by `init`).
- **Syscalls:** `fork()`, `exec()`, `wait()/waitpid()`, `exit()`, `kill()`, `getpid()`, `getppid()`.
- `fork()`: returns 0 to child, PID to parent.

Threads

- **Threads:** Execution units within a process.
- **Multithreading:** Process with multiple threads.
 - **Shared:** Memory space, FDT, directory, signal handlers.
 - **Unique:** Stack, registers (PC, SP), TID.
- **Thread Group:** Threads sharing a PID (TGID). `getpid()` → TGID, `gettid()` → TID.
- `exit()` terminates the entire thread group.
- **Pthreads:** `pthread_create()`, `pthread_join()`, `pthread_exit()`.
- **User Level Threads (ULTs):** Managed by user-library, no true parallelism.
- **Kernel Threads:** OS-level tasks, no memory space switch.

Signals

- **Signal:** Asynchronous notification.
- Wakes a waiting process.
- **Handling:** Custom handler (`sigaction()`), ignore (`SIG_IGN`), default (`SIG_DFL`).
- **Masking:** Block signals (`sigprocmask()`), becomes "pending".
- **Unblockable Signals:** `SIGKILL`, `SIGSTOP`.
- **Common Signals:** `SIGSEGV`, `SIGILL`, `SIGCHLD`, `SIGINT`, Signal 0.

IPC & System Calls

- `clone()`: Creates processes/threads. Flags: `CLONE_VM`, `CLONE_FILES`, `CLONE_SIGHAND`, `CLONE_THREAD`.
- **Copy-on-Write (COW):** `fork()` optimization, duplicates page on write.
- **Pipe:** Unidirectional buffer (`int fd[2]`).

- **FIFO (Named Pipe):** Filesystem pipe (`mkfifo()`) for unrelated processes.

Boot & Modules

- **Booting Sequence:** BIOS → MBR → Boot Loader (GRUB) → Kernel (`/sbin/init`).
- **Kernel Modules:** Dynamically loadable code (`init_module()`, `cleanup_module()`).
- **Device Files:** In `/dev`. **Major** (driver), **minor** (instance).
 - **Character:** Stream-based.
 - **Block:** Addressable.
 - **Pseudo:** Virtual (`/dev/null`).
- `file_operations` struct: Driver function pointers, registered with `register_chrdev()`.

CPU Scheduling

Concepts

- **Time Sharing:** CPU virtualization.
- **Process Types:** I/O-bound (latency-sensitive), CPU-bound.
- **Preemption:** Forcible process stop, via **Quantum**.
- **Scheduling Metrics:**
 - $T_{wait} = T_{start} - T_{submit}$.
 - $T_{resp} = T_{end} - T_{submit}$.
 - Slowdown: $1 + \frac{T_{wait}}{T_{run}}$.

Batch (Non-Preemptive) Schedulers

- **FCFS:** First-Come, First-Served.
- **SJF:** Shortest Job First.
- **EASY:** FCFS with backfilling.

Preemptive Schedulers

- **Round Robin (RR):** Quantum-based circular queue.
- **SRTF:** Shortest Remaining Time First.
- **Selfish RR:** Two-level queue with aging.
- **Gang Scheduling:** RR for parallel jobs.

Linux Schedulers

- **Linux ≤ 2.4 ($O(n)$):**
 - **HZ:** Time interrupts per second.
 - **Tick:** Time between time interrupts.
 - **Policies:** `SCHED_RR`, `SCHED_FIFO`, `SCHED_OTHER`.
 - **Epoch:** Cycle for all runnable tasks.
 - **Priority:** Static (`nice`) + Dynamic (`counter`).
 - New epoch counter: $C_{new} = \lfloor \frac{C_{old}}{2} \rfloor + P_{static}$.
 - `schedule()` finds best goodness().
- **CFS ($O(\log N)$):**
 - **Goal:** Fair CPU time via equal `vruntime`.
 - **vruntime:** Task with lowest runs.
 - $\Delta vruntime = \Delta T_{actual} \times \frac{W_{ideal}}{W_{task}}$.
 - **Data Structure:** Red-black tree ordered by `vruntime`.

Context Switch

- **Context:** Process state.
 - **Overhead:** **Direct** (save/load state), **Indirect** (cache pollution).
 - **Types:** **Forced** (interrupt), **Initiated** (syscall).
 - **Mechanism:** State saved to kernel stack. **TSS** updated.
 - **Flow:** `schedule()` → `context_switch()` → `_switch_to_asm()` → `_switch_to()`.
-

Synchronization

Concepts

- **Race Condition:** Timing-dependent outcome.
- **Critical Section:** Requires mutual exclusion.
- **Atomicity:** All-or-nothing operation.
- **Memory Consistency:** Order of memory ops, enforced by **memory fence**.
- **Amdahl's Law:**

$$Speedup = \frac{1}{s + \frac{1-s}{n}} \leq \frac{1}{s}$$

Mechanisms

- **Lock:** `acquire()` and `release()`.
- **Spinlock:** Busy-waits for lock.
- **Mutex:** May block (sleep) if lock is unavailable.
- **Semaphore:** Counter with wait queue. `wait()` (P), `signal()` (V).
- **Condition Variable:** Wait for condition inside critical section. `cond_wait`, `cond_signal`.

Deadlocks

- **Deadlock:** Processes blocked waiting for each other's resources.
 - **Livelock:** State changes, no progress.
 - **Starvation:** Process perpetually denied resources.
 - **Four Necessary Conditions:** Mutual Exclusion, Hold & Wait, No Preemption, Circular Wait.
 - **Handling:**
 - **Prevention:** Violate one condition (e.g., lock ordering).
 - **Avoidance:** **Banker's Algorithm**.
 - **Detection & Recovery:** Find cycle and kill process.
-

Networking

- **Protocol Stack:** L5-App, L4-Transport, L3-Network, L2-Link.
- **TCP:** Reliable, connection-oriented, stream-based.
- **UDP:** Unreliable, connectionless, datagram-based.
- **IP Address:** Logical host address (network part + host part).
- **Socket:** Communication endpoint (FD). 5-tuple identifies connection.
- **TCP Flow:** Server: `socket` → `bind` → `listen` → `accept`. Client: `socket` → `connect`.
- **Ethernet (L2):** LAN technology, uses **MAC addresses**.

- **Key Protocols:** **ARP** (IP → MAC), **DHCP** (dynamic IP), **NAT** (private ↔ public IP).
-

Virtual Memory

Concepts

- **Virtual Memory (VM):** Private, contiguous address space abstraction.
- **Page** (virtual, 4KB) & **Frame** (physical).
- **Address Translation:** MMU translates VA → PA via page tables.
- **Page Table:** Maps virtual pages to physical frames.
- **PTE (Page Table Entry):** Contains frame number and bits: Present, Dirty, Accessed, R/W, U/S.
- **Page Fault:** Trap on access to unmapped page.
 - **Major Fault:** Fetch from disk.
 - **Minor Fault:** Page in memory, mapping missing.
- **TLB:** Hardware cache for VA → PA translations.

Paging & Swapping

- **On-Demand Paging:** Load pages on first access.
- **Swap Area:** Disk space for evicted pages.
- **Paging out:** Copying page from DRAM to disk.
- **kswapd:** Kernel thread to reclaim page frames.
- **Thrashing:** Excessive paging, little progress.

Page Replacement Algorithms

- **Goal:** Choose victim page to evict.
- **Belady's Optimal:** Replace page used furthest in future.
- **LRU:** Replace least recently used page.
- **Clock Algorithm:** LRU approximation with reference bit.
- **Linux PFRA:** Two lists (**active**, **inactive**) to approximate LRU.

Linux Implementation

- **Multi-Level Page Tables:** 4-level on x86-64. **CR3** points to top level.
 - **mm_struct:** Process memory descriptor.
 - **vm_area_struct (VMA):** Contiguous memory region.
 - **mmap():** Creates a new VMA.
 - **Copy-on-Write (COW):** On `fork()`, share pages, copy on write.
-

Storage & Filesystems

Storage Devices

- **HDD:** Access time = seek time + rotational latency.
- **SSD:** Flash memory, low random access latency.
- **DMA:** Device-to-memory transfer without CPU.

RAID

- **RAID 0 (Striping):** Best performance, no redundancy.
- **RAID 1 (Mirroring):** Full redundancy.
- **RAID 4 (Parity Disk):** Write bottleneck on parity disk.

- **RAID 5 (Distributed Parity)**: No parity disk bottleneck.
- **RAID 6 (Dual Parity)**: Protects against two disk failures.

Filesystem Concepts

- **File**: Logical unit of information.
- **inode**: File metadata structure.
- **dirent**: Maps filename to inode number.
- **Links**: **Hard Link** (same inode), **Symbolic Link** (path

string).

- **File Descriptor (FD)**: Integer handle for an open file, index into **FDT**.
- FDT entry → system-wide **File Object** → **inode**.

Filesystem Implementation

- **Layout**: Superblock, bitmaps, inode table, data blocks.
- **Allocation**: **Multi-Level Index**, **Extents**, **FAT**.
- **Journaling**: Transaction log for consistency.
- **Path Resolution**: Traverse directories to find inode.