

Operating Systems (02340123) Summary - Spring 2025

Razi & Yara

July 19, 2025

Contents

I	Lectures & Tutorials	3
	Topic 1: Introduction	4
	Topic 2: Processes & Signals	6
	Topic 3: IPC - Inter-Process Communication	12
	Topic 4: Scheduling	15
	Topic 5: Synchronization (& Deadlocks)	25
	Topic 6: Networking	34
	Topic 7: Virtual Memory	42
II	Overall Summary	43
III	Highlights and Notes	44

Part I

Lectures & Tutorials

Topic 1: Introduction

Operating System (OS) An Operating System's job is:

- Coordinate the execution of all SW, mainly user apps.
- Provide various common services needed by users & apps.
- An OS of a physical server controls its physical devices, e.g. CPU, memory, disks, etc.
- An OS of a virtual server only *believes* it does. There's another OS underneath, called **hypervisor** which fakes it.

Using an OS allows us to take advantage of "**virtualization**":

- **Server Consolidation:** Run multiple servers on one physical server. This allows for better resource utilization, smaller spaces, and less power consumption.
- **Disentangling SW from HW:** allows for backing up/restoring, live migration, and HW upgrade. This gives us the advantage of easier provisioning of new (virtual) servers = "virtual machines", and easier OS-level development and testing.

Most importantly, an OS is **reactive**, "event-driven" system, which means it waits for events to happen and then reacts to them. This is in contrast to typical programs which run from start to end without waiting for external events to occur to invoke them.

	Typical Programs	OS
What does it typically do?	Get some input, do some processing, produce output, terminate	Waits & reacts to "events"
Structure	Has a main function, which is (more or less) the entry point	No main ; multiple entry points, one per event
Termination	End of main	Power shutdown
Typical goal	~ Finish as soon as possible	Handle events as quickly as possible \Rightarrow more time for apps to run

Event Synchronisation OS events can be classified into two:

- **Asynchronous interrupts:** keyboard, mouse, network, disk, etc. These are events that can happen at any time and the OS must be ready to handle them.
- **Synchronous:** system calls, divide by zero, page faults, etc. These are events that happen as a result of the program's execution and the OS must handle them immediately.

Multiplexing Multiplexing is the ability of an OS to share a single resource (e.g. CPU, memory, disk) among multiple processes or threads. This allows for better utilization of resources and enables multiple applications to run concurrently. *Multiprogramming means multiplexing the CPU recourse.

Notable services provided by an OS:

1. **Isolation:** Allow multiple processes to coexist using the same resources without stepping on each other's toes. Usually achieved by multiplexing the CPU, memory, and other resource done by the OS. However, some physical resources know how to multiplex themselves, e.g. network cards, sometimes called "*self-virtualizing devices*".
2. **Abstraction:** Provides convenience & portability by:
 - offering more meaningful, higher-level interfaces
 - hiding HW details, making interaction with HW easier.

Topic 2: Processes & Signals

Processes

Each process is an instance of a program in execution, which includes:

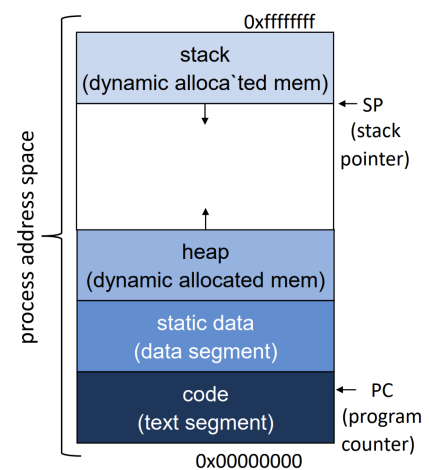
- **Program code:** The actual code of the program.
- **Process state:** The current state of the process, including the program counter, registers, and memory management information.
- **Process control block (PCB):** A data structure used by the OS to manage the process, containing information such as process ID, process state, CPU registers, memory management information, and I/O status information.

A process doesn't have direct access to its PCB, it is managed by the OS (kernel space), i.e. needs privilege level 0 (kernel mode) to access it.

Each PCB contains: `real_parent`, `parent`, `children`, `siblings`,...

Each process has a `task_struct` current pointer to its PCB.

- **Process ID (PID):** A unique identifier assigned to each process by the OS. The PID is used by the OS to manage the process and is used in system calls to refer to the process.



Process States A process can be in one of the following states:

- **Running:** The process is currently being executed by the CPU.
- **Ready:** The process is ready to be executed but is waiting for the CPU to become available.

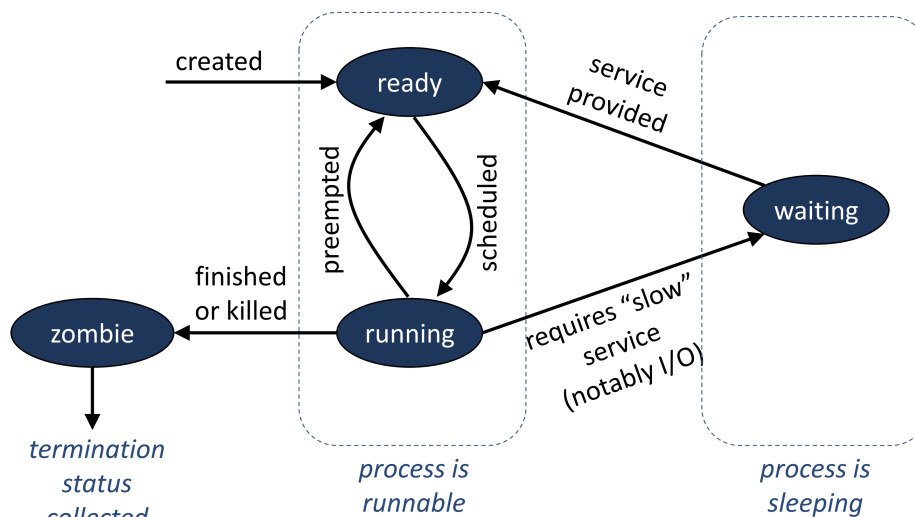


Figure 1: Process States

- **Waiting:** The process is waiting for an event to occur, such as I/O completion or a signal.
- **Zombie:** The process has terminated but its PCB is still in the system, waiting for the parent process to read its exit status. In this state, the process has released almost all of its resources, but *the PCB is still in the system*.

As we saw in "ATAM", each process can only access a certain set of utilities and functions, those who require privilege level 3 (user mode). So to access the OS services, a process must use **system calls** which are functions provided by the OS that allow processes to request services from the OS. System calls are typically implemented in the OS kernel and provide a controlled interface for processes to interact with the OS.

Each *syscall*, in case of an error, will change the **errno** variable to indicate the error type. The **errno** variable is a global variable that is set by system calls and some library functions in the event of an error to indicate what went wrong. It is defined in the header file **errno.h**. **Note:** **errno** is not reset to 0 after a successful syscall, so it must be checked immediately after the syscall, and be reset before usage if need be (if there is not any other way to make sure there is an error indeed).

i.e. `errno = <Number of Last Syscall Error>;`

As noted above, each process must be `wait()`ed for by its parent process to be able to release its PCB and resources. This is done by the `wait()` syscall, which suspends the calling process until one of its children terminates. In case *the parent process terminates before the child*, the child process *becomes an orphan process and is adopted by the init process (PID 1)*, which will then wait for it to terminate and release its resources.

Process Management The OS offers various system calls to manage processes, including: (More details in the functions reference)

- **fork()**: Creates a new process by duplicating the calling process. The new process is called the child process, and the calling process is called the parent process.
- **exec()**: Replaces the current process image with a new process image, effectively running a different program in the same process.
- **wait()**: Suspends the calling process until one of its children terminates.
- **exit()**: Terminates the calling process and releases its resources.
- **getpid()**: Returns the process ID of the calling process.
- **getppid()**: Returns the process ID of the parent process.
- **kill()**: Sends a signal to a process, which can be used to terminate or suspend the process.

Parent Vs. Real Parent Process The real parent process is the one that created the current process using **fork()**, or the one that adopted it in case the real parent terminated before the child.

The parent process is the one *tracing* the current process, e.g. using **ptrace()**. The parent process is the one that will receive signals from the current process, e.g. **SIGCHLD** when the current process terminates.

In most cases, the parent process is the real parent process, but it can be different in some cases, e.g. when a process is being traced by a debugger.

Daemon Processes A daemon process is a background process not controlled by the user. To run a process as a daemon use **nohup <command> &**.

Daemon names usually end with the letter "d", e.g. **sshd** (SSH daemon), **httpd** (HTTP daemon), etc.

Signals

Signals Signals are "notifications" sent to a process to asynchronously notify it that some event has occurred.

* Receiving a signal only occurs when returning from kernel mode, which in turn invokes the corresponding signal handler.

** Default signal handling actions: Either die or ignore

*** In case of several signals from different types, they will be handled by the order of their definition in the signals register.

Each signal has a name, a number, and a default action. All but 3 signals can be blocked, i.e. ignored until the process is ready to handle them. The 3 signals that cannot be blocked are:

- **SIGKILL**: Used to forcefully terminate a process. (Process becomes a zombie)
- **SIGSTOP**: Used to suspend the receiving process. (Make it sleep) The signal is sent when the user presses Ctrl+Z in the terminal. Note: In truth Ctrl+Z sends the SIGTSTP signal, however, we don't learn about the differences between the two signals in this course.
- **SIGCONT**: Used to resume a suspended process, usually sent after a **SIGSTOP** signal. The handler for this signal can be customized but it **will always** resume the process.

SIGSTOP and SIGCONT are useful for debugging purposes, allowing the user to pause and resume the execution of a process.

Signal Handling A process can define a custom signal handler for a specific signal using the `signal()` or `sigaction()` preferred system calls. To ignore a signal, the process can set its handler to `SIG_IGN`. To restore the default action for a signal, the process can set its handler to `SIG_DFLT`.

Signal Masking A process can block signals using the `sigprocmask()` system call, which allows the process to specify a set of signals to block. This allows the process to overcome *Race Conditions* resulted from the asynchronous nature of signals.

This is achieved by maintaining a set of currently blocked signals & a set of masked signals which is saved in the [PCB](#).

Blocked Signals = mask array.

Pending Signals = signals that were sent to the process while it was blocked, and will be handled when the process unblocks them.

Common Signals the following are some of the most common signals:

1. **SIGSEGV, SIGBUS, SIGILL, SIGFPE**: These are driven by the associated (HW) **interrupts** - The OS gets the associated interrupt, then the OS interrupt handler sees to it that the misbehaving process gets the associated signal, lastly the signal handler is invoked.
 - **SIGSEGV**: Segmentation violation (illegal memory reference, e.g., outside an array).
 - **SIGBUS**: Dereference invalid address (null/misaligned, assume it's like SEGV).
 - **SIGILL**: Illegal instruction (trying to invoke privileged instruction).
 - **SIGFPE**: Floating-point exception (despite the name, *all* arithmetic errors, not just floating point. e.g., division by zero).
2. **SIGCHLD**: Parent (and then real parent) get it whenever `fork()`ed child terminates or is **SIGSTOP**-ed.
3. **SIGALRM**: Get a signal after some specified time, can be set using the `alarm()` & `setitimer()` system calls.
4. **SIGTRAP**: When debugging/single-stepping a process, the debugger can set a breakpoint in the code, which will cause the process to receive a **SIGTRAP** signal when it reaches that point.
5. **SIGUSR1, SIGUSR2**: User-defined signals, user can decide the meaning of these signals and their handlers.
6. **SIGPIPE**: Write to pipe with no readers.
7. **SIGINT**: Sent when the user presses Ctrl+C in the terminal. The default action is to terminate the process, but it can be customized.
8. **SIGXCPU**: Delivered when a process used up more CPU than its soft-limit allows: soft-/hard limits are set using the `setrlimit()` system call. Soft-limits warn the process its about to exceed the hard-limit, Exceeding the hard-limit will cause **SIGKILL** to be sent to the process.
9. **SIGIO**: Can configure file descriptors such that a signal will be delivered whenever some I/O is ready.

Typically makes sense when also configuring the file descriptor to be *non-blocking*, e.g., when `read()`ing from a non-blocking file descriptor, the system call immediately returns to user if there's currently nothing to read. In this case, `errno` will be set to `EAGAIN=EWOULDBLOCK`.
10. **Signal 0**: The null signal, which is used to test for the existence of a process.

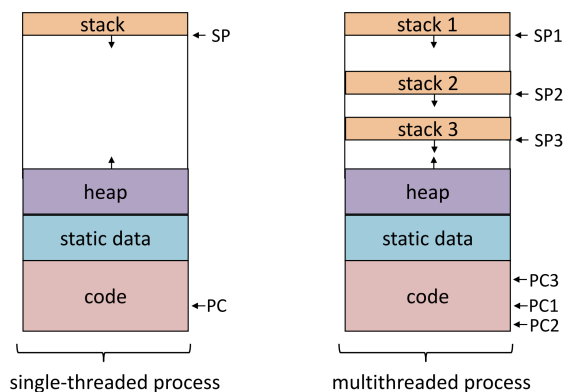
Signals Vs. Interrupts

	interrupts	signals
Who triggers them? Who defines their meaning?	Hardware: CPU cores (sync) & other devices (async)	Software (OS), HW is unaware
Who handles them? Who (un)blocks them?	OS	processes
When do they occur?	Both synchronously & asynchronously	Likewise, but, technically, invoked when returning from kernel to user

Topic 3: IPC - Inter-Process Communication

Threads & IPC

Multiprocessing Using several CPU cores (=processors) for running a single job to solve a single "problem".



Threads A process can have multiple threads, which share (nearly) everything, including the process's memory space, file descriptors, heap, static data, code segment and more.

However, each thread has its own stack, registers, and program counter. But they can still access each other's stack since they share the same memory space.

Note that *Global/Static Variables & Dynamically Allocated Memory* are shared between threads since they are stored in the Static Data Segment and the Heap, respectively, which are shared between all threads of a process. However, *Local Variables* are not shared between threads since they are stored in the stack, which is unique to each thread.

	Unique to Process	Unique to pthread
Registers (notably PC)	Y	Y
Execution stack	Y	Y
Memory address space	Y	N
Open files	Y	N
Per-open-file position (=offset)	Y	N
Working directory	Y	N
User/group credentials	Y	N
Signal handling	Y	N

OpenMP Open Multi-Processing (OpenMP) consists of a set of compiler *'pragma'* directives that allows the compiler to generate multi-threaded code for parallel execution.

```

1      #pragma omp parallel for
2      for (i = 0; i < N; i++) {
3          arr[i] = 2*i;
4      }
```

Pthreads POSIX threads (pthreads) is a standard for multi-threading in C/C++. It provides a set of functions to create and manage threads, as well as to synchronize them. (More details in the functions reference)

File Descriptors A non-negative integer representing an I/O "channel" on some device. File descriptors are saved in the process's PCB inside the FDT (File Descriptor Table), which is an array of pointers to *file objects*, each representing an open file or device. So in fact, an FD is an index to a kernel array of channels.

Processes **don't** share PCB nor FDT but they **can share file descriptors**, i.e. two processes can have the same file descriptor pointing to the same file object, which allows them to share the same open file or device. Upon forking, the child process inherits a copy of the parent's FDT. **Note:** Threads **do** share the same FDT.

Pipes A pipe is a unidirectional communication channel between two processes, allowing one process to send data to another. Pipes are a pair of two file descriptors `int pipe_fd[2]`. Each integer is a handle to a kernel communication object, which is a buffer that holds the data being sent between the two processes.

- `pipe_fd[0]` = read side of the communication channel.
- `pipe_fd[1]` = write side of the communication channel.
- Everything written via `pipe_fd[1]` can be read via `pipe_fd[0]`.
- **Blocking:** If the read side is empty, the read operation will block until data is written to the write side. If the write side is full, the write operation will block until space is available.
- **SIGPIPE:** Writing to a pipe whose read end is `close()`d will result in a SIGPIPE signal.

Multi-tasking	Multi-programming	Multi-processing
Having multiple processes <i>time slice</i> on the same CPU core.	Having <i>multiple jobs</i> in the system (either on the same core or on different cores). i.e. the existence of multiple processes in the system, regardless of whether they are running on the same core or not.	Using <i>multiple processors (CPU cores)</i> for the same <i>job</i> in parallel. i.e. creating multiple threads to run on different cores for the same process.

Intro. Context Switching & Caching

Context Switching is the process of saving the state of a currently running process and restoring the state of another process to allow it to run. This is done by the OS kernel and is necessary for multitasking, allowing multiple processes to share the CPU.

- **Context** = the state of a process, including its registers, program counter, stack pointer, and memory management information.
- **Context Switch** = the process of saving the context of the currently running process and restoring the context of another process.

Context Switch Overhead consists of two components:

- **Direct Overhead:** The measurable time it takes to perform the context switch, which includes saving the current process's state and restoring the next process's state.
- **Indirect Overhead:** The time it takes for the CPU to cache the new process's data, which can be significant if the new process's data is not already in the CPU cache.

Caching The CPU cache is a small, fast memory that stores frequently accessed data to speed up access times. The Cache allows the CPU access to a fast memory that is closer to the CPU than the main memory (DRAM), and a larger one than the CPU registers. The cache works because of the *Principle of Locality*:

- **Temporal Locality:** If at one point in time a particular memory location is referenced, then it is likely that the same location will be referenced again soon.
- **Spatial Locality:** If a particular memory location is referenced at a particular time, then it is likely that nearby memory locations will be referenced soon.

Copy-on-Write (COW) The `fork()` system call creates a copy of the address space of the parent, but:

- It only creates a *logical* copy.
- There is no physical duplication of the memory pages, until we really need to write to them (from either process). And even then, only the page that is being written to is duplicated. (More details in the virtual memory lectures)

User Level Threads (ULTs) are threads that are managed by the user-level library, rather than the OS kernel. ULTs are not visible to the OS, which means that the OS does not know about them and does not schedule them. This allows for faster context switching between ULTs, but it also means that the OS cannot take advantage of multiple CPU cores to run ULTs in parallel.

The usage of ULTs is mainly for concurrent programming, where we want multiple multiple tasks to progress in parallel without the overhead of kernel-level threads.

Topic 4: Scheduling

Batch (Non-Preemptive) Scheduling

Supercomputers Supercomputers are comprised of multiple nodes, each with multiple CPU cores, and are used for running large-scale computations. Users submit **batch jobs** to the supercomputer with a specified time limit and size. Terminology:

- **Size** = the number of CPU cores to use for the job. Jobs are said to be *wide/big* or *narrow/small*.
- **Runtime** = the time limit for the job to run. Jobs are said to be *short* or *long*.

Metrics for Performance Evaluation When evaluating the performance of a batch scheduling algorithm, we use the following metrics:

- **Average Wait Time:** The wait time of a job is the interval between the time the job is submitted to the time the job starts to run. i.e.

$$\text{waitTime} = \text{startTime} - \text{submitTime}$$

- **Average Response Time:** The response time of a job is the interval between the time the job is submitted to the time the job is terminated. i.e.

$$\text{responseTime} = \text{terminateTime} - \text{submitTime}$$


Note: Average wait time and response time differ only by a constant factor, which is the job's average runtime.

- **The Slowdown / Expansion Factor:** The ratio between a job's response time and its runtime. i.e.

$$\begin{aligned}\text{slowdown} &= \frac{\text{responseTime}}{\text{runtime}} \\ &= \frac{(\text{waitTime} + \text{runtime})}{\text{runtime}} \\ &= 1 + \frac{\text{waitTime}}{\text{runtime}}\end{aligned}$$

- **Utilization:** The percentage of time the resource (CPU) is busy.
- **Throughput:** How much work is done in one time unit.

Batch Scheduling Algorithms




First-Come, First-Served (FCFS) Jobs are scheduled by their arrival time. If there are enough free cores, a newly arriving job starts to run immediately, otherwise it wait, sorted by arrival time, until enough cores are available.

- **Pros:** Simple to implement (FIFO Queue), fair to all jobs.
- **Cons:** Creates fragmentation, small/short jobs might wait a long time.

EASY Scheduling (= FCFS + backfilling) Back-filling optimization: A short waiting job can jump over the head of the wait queue (i.e. start at an earlier time) provided that it **doesn't delay** the job head of the FCFS queue.

The algorithm: whenever a job arrives or terminates, try to start the job head of the FCFS wait queue. Then, iterate over the rest of the waiting jobs (in FCFS order) and try to backfill them.

- **Pros:** Better utilization (less fragmentation), short jobs have a better chance of running sooner.
- **Cons:** Must know runtimes in advance.



Shortest Job First (SJF) Instead of ordering jobs by their arrival time, we order them by **their (typically estimated) runtimes**.

- **Pros:** **Optimal** in terms of performance (min. avg. wait time).
- **Cons:** Unfair as it may cause **starvation** of long jobs. starvation = can theoretically wait forever

Convoy Effect Slowing down all (possibly short) processes due to currently servicing a very long process. *FCFS* suffers from this effect, as does *EASY* scheduling but to a lesser

extent. *SJF without assumptions* also suffers from this effect but to an even lesser extent.

Optimality of SJF As mentioned above, SJF is optimal in terms of performance, in particular, it minimizes the average wait time. **Claim:**

Given:

1. A 1-core system where all jobs are serial.
2. All process arrive together.
3. Their runtimes are known in advance.

Then: *The average wait time of SJF is equal to or less than the average wait time of any other batch scheduling order.*

Fairer Variants of SJF Motivation: disallow job starvation.

- **Shortest-Job Backfilled First (SJBFBF):** Exactly like EASY in terms of servicing the head of the wait queue in FCFS order (and not allowing anyone to delay it), but the *backfilling* traversal is done in SJF order, i.e. the next job to be backfilled is the one with the shortest estimated runtime.
- **Largest eXpansion Factor (LXF):** LXF is similar to EASY, but instead of ordering the wait queue in FCFS, it orders jobs based on their current slowdown (**greater slowdown = higher priority**).

On every job arrival or termination, the expansion factors are recalculated and the wait queue is resorted, so that the job with the largest expansion factor is always at the head of the queue. Both the head of the queue and the backfilled jobs are serviced in LXF order.

Preemptive Schedulers

Preemption is the act of suspending one job (process) in favor of another even though it is not finished yet.

Why do we need preemption? Preemption is necessary for **responsiveness** and when the runtime of jobs vary or unknown in advance.

Quantum is the maximum amount of time a process is allowed to run before it is pre-empted. Quantum is typically milliseconds to 10s of milliseconds, and is *often set per-process*. Usually a CPU-bound process gets long quanta while an I/O-bound process gets short quanta with higher priority.

Performance Metrics for Preemptive Schedulers we use the following:

- **Average Wait Time:** As before, the wait time of a job is the interval between the time the job **is submitted** to the time the job **starts to run**. *Note* that the wait time **does not include** the preemption wait times (i.e. the time the job is waiting for its turn to run again after being preempted).
- **Response Time (=Turnaround Time):** Like before, the response time is the time from process submission to process completion. *Note* that with preemption we get:

$$\text{responseTime} \geq \text{waitTime} + \text{runtime}$$

due to the preemption wait times and **Context Switches** cost.

- **Overhead:** How long a context switch takes, and how often context switches happen. (Should be minimized)
- **Utilization & Throughput:** Same as before, but we may want to account for context switch overhead.
- **Makespan Time:** The time it takes for the system to finish all jobs in the system.

Round Robin (RR) Scheduling Processes are arranged in a cyclic read-queue. The algorithm works in the following steps:

1. The head process runs until its quantum is exhausted.
2. The head process is then preempted (suspended) and moved to the tail of the queue.
3. The scheduler resumes the next process in the circular list.
4. When we've cycled through all processes in the run-list (and we reach the head process again), we say that the current **"epoch"** is over, and the next epoch begins.

Note: RR Requires a timer interrupt; Typically it's a periodic interrupt (fires every few milliseconds) and upon receiving the interrupt, the OS checks if its time to preempt. **Note 2:** For small enough quantum, it's like everyone of the N processes advances in $1/N$ of the speed of the core called sometimes virtual time. With huge quantum infinity, RR becomes FCFS.

Gang Scheduling Think of it as RR for *parallel* systems. It works as follows:

- Time is divided to slots (seconds or minutes).
- Every job has a **"native"** time slot.
- Algorithm attempts to fill holes in time slots by assigning to them jobs from other native slots (called **"alternative"** slots).
- Algorithm attempts to minimize slot number using **slot unification** when possible.
- Rarely used in practice, if lots of memory is swapped out upon context switch, context switch overhead is too high.

Detailed explanation about alternative slots and slot unification:

Alternative Slots To further enhance resource utilization, the gang scheduling algorithm attempts to fill holes in time slots. A job has a "native" time slot where it is primarily scheduled to run. However, if there are idle processors in other time slots (termed "alternative" slots), the scheduler can assign jobs from their native slots to run in these otherwise unused resources. This allows jobs to get more processing time without interfering with the primary scheduling of other jobs.

Slot Unification This technique aims to minimize the total number of active time slots. If, after some jobs have completed, the remaining jobs in two different time slots can fit into a single slot without conflict, the scheduler can merge them. By unifying slots, the overall cycle time of the matrix is reduced, which can lead to shorter job turnaround times and improved system throughput.

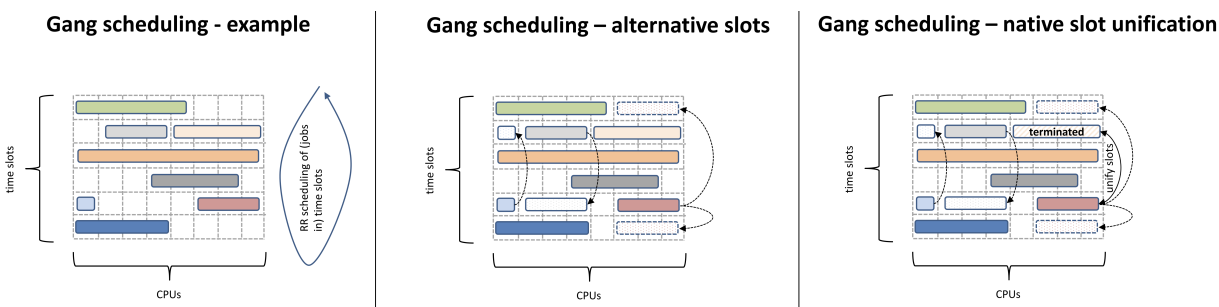


Figure 2: Gang Scheduling

Batching vs. Preemption **Assume:** (1) A single core system (2) All jobs arrive together (3) Context switch price = 0. **Then:** There exists a non-preemptive algorithm such that:

$$\text{avgResponseTime}(\text{non-preemptive}) \leq \text{avgResponseTime}(\text{preemptive})$$

As a result, SJF is also optimal relative to preemptive scheduling (if our assumptions hold).

$$\text{avgResponseTime}(\text{SJF}) \leq \text{avgResponseTime}(\text{batch or preemptive scheduler})$$

Connection Between RR & SJF **Assume:** (1) A single core system (2) All jobs arrive together (and only use CPU, no I/O) (3) Quantum is uniform + **no** context switch overhead. **Then:**

$$\text{avgResponse(RR)} \leq \text{avgResponse(SJF)}$$

i.e. RR at best is as good as SJF, and at worst is $2X$ slower.

Shortest Remaining Time First (SRTF) If we remove the assumption that all jobs arrive together, then SJF is not optimal anymore, as it may cause a convoy effect. **SRTF** is a preemptive version of SJF, where the scheduler always runs the job with the shortest remaining time. i.e. when a new job arrives or an old one finishes, the scheduler runs the job with the shortest remaining time, even if it means preempting the currently running job. This is optimal in terms of average wait time, but it can cause starvation for long jobs.

Selfish RR Selfish RR is a variant of RR where only "old enough" processes are running in the run-queue. i.e.

- New processes wait in a FIFO queue, not yet scheduled.
- Older processes scheduled using RR.
- New processes are scheduled when either of the following happens:
 1. The run-queue is empty (i.e. no ready-to-run "old" processes exist).
 2. "Aging" is being applied to new processes (a per-process counter increases over time); When the counter passes a certain threshold, the "new" process becomes "old" and is transferred to the RR queue.
- **Fast Aging** = Algorithm resembles RR, **Slow Aging** = Algorithm resembles FCFS.

General Purpose Schedulers (Priority-Based & Preemptive)

Scheduling using priority Every process is assigned a priority. The priority reflects how "important" it is in that time instance, and it changes dynamically over time. Process with higher priority are favored, as they are scheduled before lower priority processes. The concept of priority can be also applied in batch scheduling. e.g. SJF:priority = runtime (smaller runtime = higher priority), FCFS:priority = arrival time (earlier arrival = higher priority), etc.

Negative Feedback Principle The schedulers of all general-purpose OSes employ a negative feedback policy: Running reduces priority to run more & Not running increases priority to run. As a result, I/O-bound processes (that seldom use the CPU) get higher priority than CPU-bound processes (that use the CPU a lot). This ensures that I/O-bound processes are responsive.

Multi-Level Priority Queue A multi-level priority queue consists of several RR queues, each associated with a priority. processes migrate between the queues so they have a dynamic priority, "important" processes move up (e.g. I/O-bound) and "unimportant" processes move down (e.g. CPU-bound). Priority is greatly affected by the **negative feedback principle**, i.e. by CPU consumption:

- I/O-bound \iff move up
- CPU-bound \iff move down

Some schedulers allocate short quanta to higher priority queues, and some don't or even do the opposite.

Linux \leq 2.4 Scheduler

All of the definitions below are relative only to the Linux \leq 2.4 scheduler, which is a preemptive priority-based scheduler that uses a multi-level priority queue. **Note:** A task is either a process or a thread.

Standard POSIX Scheduling Policies POSIX dictates that each task is associated with one of three scheduling policies:

- "Realtime" Policies:
 1. SCHED_RR (round-robin)
 2. SCHED_FIFO (first-in, first-out)
- The default policy:
 3. SCHED_OTHER (the default time-sharing policy)

POSIX defines the meaning of SCHED_OTHER (aka SCHED_NORMAL in Linux) is decided by the OS. Typically employs some multi-level priority queue with the negative feedback loop.

Realtime tasks are **always favored** by the scheduler. Only an admin can set a task to run with a *realtime* policy, the users can set a policy using the `sched.schedscheduler` syscall.

Epoch As mentioned before, every runnable task gets allocated a quantum, which is the CPU time the task is allowed to consume before it's stopped by the OS. Whenever all quanta of all *runnable* tasks become zero a **new epoch** begins. In this case we allocate additional running time to *all tasks* (runnable or not).

Definitions

- **Task's Priority:** Every task is associated with an integer, the higher the value the higher the priority to run. Every task has a *static* priority and a *dynamic* priority.

- **Static Priority:** A fixed value that indirectly determines the maximal quantum for the task. Fixed unless the user invokes `nice()` or `sched_setscheduler` syscalls.
- **Dynamic Priority:** Is the both the *remaining time* for the task and its *current priority*. The dynamic priority decreases over time (while running), if it reaches zero the task is preempted until the next epoch. The dynamic priority is *reset* to the static priority at the beginning of each epoch.
- **HZ:** Linux gets a timer interrupt *HZ* times per second, i.e. it gets a timer interrupt every $\frac{1}{HZ}$ seconds. The default value of *HZ* is 100 for x86/Linux2.4.
- **Tick:** A tick can mean either of the following:
 - The time that elapses between two consecutive timer interrupts, i.e. $\frac{1}{HZ}$ seconds.
 - The timer interrupt itself that fires every $\frac{1}{HZ}$ seconds.

Ticks are used to determine the scheduler timing resolution, the OS measures the passage of time in ticks. The units of the *dynamic priority* are ticks.

- **task_struct:** Every task is represented by a *task_struct* object, which contains (among other things):
 1. nice static priority
 2. counter dynamic priority
 3. processor the last CPU core the task ran on
 4. need_resched boolean
 5. mm task's memory address space
- **task's nice** There are two types of nice, user nice and kernel nice. The **kernel's nice** is the *static priority* of the task, which is *between 1..40* (higher is better) and the *default is 20*. The **user's nice** is the parameter passed to the `nice()` syscall, which is *between -20...19* (*lower* is better). Values below 0 require superuser privileges.

$$\text{kernel_nice} = 20 - \text{user_nice}$$

- **task's counter** The *dynamic priority* of the task. It is calculated as follows:

- Upon task creation:

$$\begin{aligned} \text{child.counter} &= \text{parent.counter} / 2; && \text{(round down)} \\ \text{parent.counter} - &= \text{child.counter}; && \text{(round up)} \end{aligned}$$

- Upon a new epoch:

$$\begin{aligned} \text{task.counter} &= \text{task.counter} / 2 + \text{NICE_TO_TICKS}(\text{task.nice}) \\ &= \text{half of prev dynamic} + \text{convert_to_ticks}(\text{static}) \end{aligned}$$

- When running: decrement each tick by 1 ($\text{task.counter} -$) until it reaches 0.

The **NICE_TO_TICKS** function scales 20 (=DEF_PRIORITY) to number of tick compromising **50+ ms**. By default, scales 20 to 5+ ticks:

```
#define NICE_TO_TICKS(kern_nice) ((kern_nice) / 4 + 1)
```

So the quantum range is therefore: (recall that 1 tick = 10 ms)

- $(1/4 + 1 =) 1 \text{ tick} = 10 \text{ ms (min.)}$
- $(20/4 + 1 =) 6 \text{ ticks} = 60 \text{ ms (default)}$
- $(40/4 + 1 =) 11 \text{ ticks} = 110 \text{ ms (max.)}$

- **task's processor** Logical ID of CPU core upon which task has executed most recently, if task is currently running, then this is the core it is running on.
- **task's need_resched** A boolean flag that is checked by kernel just before switching back to user-mode. If set, check if there's a "better" task than the one currently running, and if so, switch to it. Can be thought of as a per-core rather than per-task flag as it is checked only for the currently running task.
- **task's mm** A pointer to the task's memory address space. (More details in the virtual memory lectures)

The scheduler is implemented in the `kernel/sched.c` file, and the task structure is defined in `include/linux/sched.h`. The scheduler is compromised of **4 main functions**:

1. `goodness(task, cpu)` - Given a task and a CPU core, returns how "desirable" it is for that CPU. We compare tasks by this value to determine which task to run next.
2. `schedule()` - Actual implementation of the scheduler. Uses `goodness` to decide which task to run on a given core.
3. `__wake_up_common(wait_queue q)` - Wakes up task(s) when waited-for event has happened. e.g. completion of I/O operation, or a signal being sent to the task.
4. `reschedule_idle(task t)` - Given a task, check whether it can be scheduled on some core. Preferably on an idle core, but if not then by preempting a less "desirable" task on a busy core. **Note:** used by the `schedule()` & `__wake_up_common()` functions.

Counter Convergence Claim: The counter value of an I/O-bound task will quickly converge to 2α geometric series. **Corollary:** By default, an I/O-bound task will have a counter of 12 ticks (=120 ms) as long as it remains I/O-bound and consumes negligible CPU time.

The Drawback The scheduler is a linear scheduler, i.e. it runs in $O(n)$ time. It was replaced by the " $O(1)$ " scheduler which then was replaced by the "*Completely Fair Scheduler*" (CFS) which is $O(\log n)$.

Topic 5: Synchronization (& Deadlocks)

Race Condition We say that the program suffers from a *race condition* if the *outcome is nondeterministic and depends on the timing of uncontrollable, un-synchronized events*

Memory (Cache) Coherency is the mechanism that ensures that any change to a shared piece of data in one cache is eventually propagated to all other caches, preventing processors from working with outdated information.

Cache Consistency it dictates how memory operations (reads and writes) from one processor can be observed by other processors in the system. If store operations are immediately seen by other cores and they can't be reordered with load operations, then consistency holds. (but it's not a necessary condition)

Cache **coherency** relates to a **single** memory location, while cache **consistency** relates to **multiple** memory locations. coherency ensures cores see writes in some order that makes sense, and is uniform across all cores. Consistency is about the order of different read and write operations across multiple memory locations, and how they are observed by different cores.

To enforce ordering for memory consistency, explicit **memory fence** (memory barrier) must be used. The fence makes all store-s that happened before the fence visible to all load-s after the fence, notably, flushes the store buffer to the memory system.

Atomicity In programming, atomicity means an operation, or a series of operations, appears to happen all at once, without any interruption from other processes. This ensures that either the entire operation completes successfully, or it has no effect at all, preventing partial or incomplete results.

Critical Section A group of operations we need to atomically execute is called a **critical section**. Atomicity will make sure other threads don't see partial results! The critical section can be a uniform code across all threads, or it can be a different.

Mutual Exclusion (mutex) is a property of a critical section that ensures that only one thread can execute it at a time and never simultaneously. Thus a critical section is a "serialization point".

Locks

Lock A lock is an abstraction that supports two operations: `acquire(lock)` and `release(lock)`. Semantically: It's a memory fence, only one thread at a time can acquire the lock and other simultaneous attempts to acquire are postponed until the lock released. Implementing a lock is nowadays done using hardware support with special instructions that ensure mutual exclusion.

Spinlock A spinlock is a lock that uses busy-waiting to acquire the lock. i.e. it repeatedly checks if the lock is available and only then acquires it. Spinlocks are typically used in low-level code where the overhead of sleeping and waking up is too high, or when the critical section is very short. If interrupt or signal handlers access the same data as the critical section, then the we must disable them, or make sure they acquire the lock too!

```
1 struct spinlock {
2     uint locked; // 0 = unlocked, 1 = locked
3 };
4
5 // The volatile keyword is crucial here. It tells the compiler that the value at addr
can be changed by external factors (like another CPU core or hardware) at any time.
This prevents the compiler from making optimizations like caching the value in a
register, forcing it to read directly from memory every time.
6 inline uint xchg(volatile uint *addr, uint newval) {
7     // atomic [oldval = *addr, swap(*addr, newval)]
8     // xchg is atomic
9     // lock adds a memory fence
10    uint oldval;
11    asm volatile("lock; xchg %0, %1"
12                : "+m" (*addr), "=r"(oldval):
13                : "1"(newval):
14                : "cc");
15    return oldval;
16 }
17
18 void acquire(struct spinlock *lock) {
19     disable_interrupts();
20     while (xchg(&lock->locked, 1) != 0) {
21         // busy-wait until the lock is available
22         enable_interrupts(); // if we want to enable
23         disable_interrupts(); // interrupts while spinning
24     }
25 }
26
27 void release(struct spinlock *lock) {
28     xchg(&lock->locked, 0); // set the lock to unlocked
29     enable_interrupts();
30 }
```

Listing 1: Spinlock Implementation

Other Atomic Operations We know of **3** atomic operations that are used to implement spinlocks:

1. **xchg(addr, newval)** - Atomically exchanges the value at **addr** with **newval** and returns the old value.
2. **test_and_set(bool *b)** - Atomically sets ***b** to *true* and returns the old value of ***b**.
3. **cas(int *p, int old_val, int new_val)** - (Compare-and-Swap) Atomically compares the value at ***p** with **old_val**, and if they are equal, sets ***p** to **new_val** and returns true. Otherwise, it returns false and does not change ***p**.

Spinning or Blocking A lock can be implemented using either **spinning** or **blocking**:

- **Spin (busy wait)** - The thread repeatedly checks if the lock is available, consuming CPU cycles while waiting. This is suitable for short critical sections where the wait time is expected to be very short. Note that the user **can not** block context switches.
- **Block (go to wait/sleep)** - The thread goes to sleep and is woken up when the lock is available. This is suitable for longer critical sections where the wait time can be significant.

Rule of thumb: If the critical section is shorter than the time a context switch takes, then use spinning, otherwise use blocking.

Unlike spinning which can be done by the user, going to sleeping is done by a request to the OS (since it involves changing process states).

Semaphores

The Basics A semaphore is usually implemented as a counter and a queue of waiting threads. A task that announces it's waiting for a resource will get the resource if it is available or will go to sleep otherwise. In case it goes to sleep it will be awakened when the resource becomes available to it.

The fields of the semaphore are as follows:

- Value (integer) (the counter):
If value is Non-negative, it indicates the number of available resources. If value is negative, it indicates the number of tasks waiting for the resource.
- A queue of waiting task:
Every task that is waiting for the resource to become available is stored in this queue. When the value is negative, $|\text{value}| = \text{queue.length}|$

The functions used to manipulate the semaphore are:

- **wait(sem)**: Decrements the value of the semaphore. If the value *after decrementing* is negative, the task goes to sleep and is added to the queue of waiting tasks. Otherwise, it continues execution.

- **signal(sem)**: Increments the value of the semaphore. If the value was negative before incrementing, it wakes up one task from the queue of waiting tasks.

Semaphore Implementation The semaphore is implemented in the kernel and **not** part of the course material. However, the following is a demonstration of the problem known as **lost wakeups**.

```
1 struct semaphore_t {
2     int value;
3     wait_queue_t wait_queue;
4     lock_t lock;
5 };
6
7 void wait(semaphore_t *sem) {
8     lock(&sem->lock);
9     sem->value--;
10    if (sem->value < 0) {
11        enqueue(self, &sem->wait_queue);
12        unlock(&sem->lock);    // This is the problem!
13        block(self);
14    } else
15        unlock(&sem->lock);
16 }
17
18 void signal(semaphore_t *sem) {
19     lock(&sem->lock);
20     sem->value++;
21     if (sem->value <= 0) {
22         p = dequeue(&sem->wait_queue); // Part of the problem too!
23         wakeup(p);
24     }
25 }
```

Listing 2: Lost Wakeups Problem

The problem with the above implementation is that we release the the lock before going to sleep but after joining the queue, so if another task calls **signal()** redbefore we go to sleep, it will wake us up (without being asleep) and we **will exit** the wait queue so *no one will be*

able to wake us up again. This is known as the **lost wakeups** problem.

Spinlock vs. Semaphore

	spinlock	semaphore
wait by granularity	spinning fine	sleeping coarse: might wait for a long time
complexity	lower	greater: typically uses lock
expressiveness	lower	greater: equivalent to lock if maximal value=1 (called “binary semaphore”); otherwise, counting how many resources are available
interface ordering	strict: must acquire before release, must release after acquire	relaxed: may signal(s) without wait(s), may wait(s) without signal(s)
interface dependence	strict: release only invoked by locker	relaxed: signal may be invoked by threads that didn't previously wait and vice versa

Amdahl's Law Amdahl's Law is a formula that gives the maximum improvement to an overall system when only part of the system is improved.

So what is the maximal expected speed when parallelizing?

- Let n be the number of threads.
- Let s be the fraction of the program that is strictly serial ($0 \leq s \leq 1$).
- Let T_n be the time it takes to run the algorithm with n threads.
- Then, optimally:

$$T_n = T_1 \times \left(s + \frac{1-s}{n} \right) \geq T_1 \times sem$$

- The speedup is defined as:

$$speedup = \frac{T_1}{T_n} = \frac{1}{s + \frac{1-s}{n}} \leq \frac{1}{s}$$

Deadlocks & Livelocks

Progress At any given time, at least one thread is able to do the critical section, i.e. No "deadlock" or "livelock" occur.

- **Deadlock** - A situation in which two or more competing actions are each waiting for the other to finish, and thus no one ever finishes. e.g. two threads waiting for each other to release a lock.
- **Livelock** - A situation in which two or more competing actions are each trying to finish, but they keep interfering with each other, and thus no one ever finishes. e.g. two threads trying to acquire a lock, but they keep releasing it and trying again.
- **Starvation** - A situation in which a thread is unable to finish its critical section because it is always preempted by other threads. e.g. a low-priority thread that is never scheduled because high-priority threads keep running.

Deadlock State We say that a system is in a deadlock state, and not a given process.

The popular definition: A set of processes is deadlocked if each process in the set is waiting for an event that only a process in the set can cause. **Our definition:** A set of processes is deadlocked if **each** process in the set is waiting for a resource held by another process **in the set** (The set can be a subset of the set of all process). **Note:** Deadlocks are tough to deal with since the system is not guaranteed to enter a deadlock if it is prone to deadlocks.

Resource types Resource **types** can be either physical (e.g. Disk/Disk Block, NIC (network controller), Keyboard, etc.) or logical (e.g. Process Table / node in a tree, etc.).

Some resources only have one instance (e.g. keyboard) and others have multiple instances (e.g. producer-consumer ring).

Therefore, **Locks** are a synchronization construct that reflect a single instance, while **Semaphores** are a synchronization construct that reflect multiple instances of a resource type. **Note:** A binary semaphore is equivalent to a mutex lock.

Resource Allocation Graph When considering resource management, it is convenient to represent the system state with a directed graph. We can represent the system in one of two ways:

- **Multi-Resource Allocation Graph:** A directed graph where: (1) "round" nodes are processes (2) "square" nodes are resources with the dots in them representing the number of instances (3) edges from process to resource represent request (4) edges from resource to process represent allocation.
- **Single Instance Resources (Simplified):** We can simplify the system with a directed graph where: (1) nodes are processes (2) edges mark dependencies between processes, i.e. an edge from process A to process B means that A is waiting for a resource held by B.




Figure 3: Resource Allocation Graph

Necessary Conditions for Deadlocks All of the following **must hold** for a deadlock to occur:

1. **Mutual Exclusion:**

At least one resource must be held in a non-shareable mode, i.e. only one process at a time can use the resource. If every resource is shareable or used by a single process, then no deadlock can occur.

2. **Hold & Wait:**

Processes holding resources are allowed to request additional resources without releasing the resources they already hold. If processes must release all resources before requesting new ones or atomically request all resources, then no deadlock can occur.

3. **Circular Wait:**

There must be a circular chain of two or more processes, each waiting for a resource held by the next process in the chain. i.e. $P(i)$ wait for resource held by $P((i+1)\%n)$. Otherwise there exists at least one process that doesn't need to wait.

4. **No Resource Preemption:**

Resources cannot be forcibly taken from a process holding them, they must be voluntarily released. If resources held can be released after some period of time, then no deadlock can occur.

Dealing with Deadlocks There are 3 main strategies to deal with deadlocks:

- **Deadlock Prevention:** Violate one of the 4 conditions necessary for deadlock to occur.
- **Deadlock Detection & Recovery:** Allow the system to enter deadlock state, but put in place mechanisms that can detect, and then recover from this situation. Typically refers to killing one or more processes in the deadlock state, using the algorithmic resource-graph problem.

- **Deadlock Avoidance:** Use a resource allocation graph to detect potential deadlocks and avoid them by not allocating resources that would lead to a deadlock. (e.g. Banker's Algorithm). Requires full knowledge of the system's resource requirements and current allocation.

Deadlock Prevention There are 4 conditions for deadlock to occur, and we can prevent deadlocks by violating one of them:

1. **No Mutual Exclusion:**

Making all resources sharable, by implementing some data structure using only hardware-supported atomic operations. These algorithms are called "lock free".

2. **No Hold & Wait:**

Each process must request all resources it needs at once, or release all resources before requesting new ones. The system either grants all resources or none and block the process. **Con:** Processes will hold on to their resources for more time than they actually need them, which limits concurrency and hence limits performance.

3. **No Circular Wait:** the most usable flexible method

Impose a strict ordering on resources, i.e. number all resources in one sequence. Processes can only request resources in increasing order, a process that wants to acquire a resource that has a lower order must first release all resources it holds that have a higher order.

4. **No "No Resource Preemption":**

When needed, choose a victim process and release its resources, allowing other processes to continue.

e.g. if there isn't enough memory, we can backup the victim's state to disk and release all its memory, or just kill the process.

Deadlock Detection & Recovery If there is only one instance of each resource type, then a deadlock can be detected by checking for cycles in the **simplified resource allocation graph**. In the general case (multiple instances of each resource type), Necessary conditions for deadlock are not sufficient, nevertheless, we can still detect deadlocks but it is *outside the scope of this course*.

After a deadlock is detected, we can recover from it by: **(1)** Killing one or more processes in the deadlock state until the deadlock is resolved, or **(2)** Forcibly preempting resources from one or more processes in the deadlock state. Finding the minimal set of processes to kill/preempt is a hard problem.

In literature, Detection and Recovery refer to the algorithmic problem of finding the minimal set of processes to kill/preempt in the resource graph.

Deadlock Avoidance The system can avoid deadlocks by not allocating resources that would lead to a deadlock. This is achieved by the **Banker's Algorithm**, which is a resource allocation and deadlock avoidance algorithm that tests for safety before allocating resources. The algorithm maintains a *safe state* and only allocates resources if it can guarantee that the system will remain in a safe state after the allocation. The algorithm requires knowledge of the maximum resource requirements of each process and the current allocation of resources.

Safe State: A system state where we are sure that all processes can be executed **in some order**, one after the other, such that each one will obtain all the resources it needs to finish. We assume once a process receives all the resources it needs, it will finish and release all its resources.

Rules of the Banker's Algorithm The Banker's Algorithm works as follows:

- There are n processes and k resource types (each resource can have multiple instances).
- Upon initialization, each process p declares the maximal number of resource-instances it will need for each type, $\text{max}[p]$.
- While running, OS maintains how many resources are currently allocated to each process, $\text{cur}[p]$.
- And how many resource instances are currently available (free), avail .
- Upon a request for resources by a process p , R , the OS allocates iff the system will stay in a safe-state after the allocation. If not, the process is blocked until the resources are available.

Note: The Banker's Algorithm is *conservative*, i.e. it will not allocate resources even if there is possibility to allocate them without deadlocking the system. The OS doesn't take chances! With Time Complexity $O(n^2)$.

```

1 // *Assume* that request R was granted to process q
2 cur[q] += R; // vector addition
3 avail -= R;
4 // Check if the system is in a safe state (i.e. all processes can finish in some order)
5 P = ... // All non-terminated processes (processes that are not finished yet)
6 while(P is *not* empty){
7     found = false // true = found a process that can finish
8     for each p in P { // find a process that can finish
9         if (max[p] - cur[p] <= avail) { // can finish? (by maxing its resources)
10             avail += cur[p] // release resources (pretend p finishes)
11             P.remove(p)
12             found = true
13         }
14     }
15     if (!found) return FAILURE
16 }
17 return SUCCESS

```

Listing 3: Banker's Algorithm

Topic 6: Networking

Protocol A protocol is a set of rules defining the *format & order* of messages sent & received, and what actions are taken upon message transmission or receipt.

Network protocol stack consists of layers




Figure 4: Network Layers Overview

In a nutshell, networks can be divided into 5 layers as follows:

- **Application Layer - L5:** The top layer, where applications communicate with each other. Examples include HTTP, FTP, SMTP, etc.
- **Transport Layer - L4:** Responsible for end-to-end communication between applications. It provides reliable or unreliable delivery of data. Examples include TCP (reliable) and UDP (unreliable).
- **Network Layer - L3:** Responsible for routing data between devices on different networks. It provides logical addressing and routing. Examples include IP (Internet Protocol).

- **Data Link Layer - L2:** Responsible for node-to-node communication within the same network. It provides physical addressing and error detection. Examples include Ethernet, Wi-Fi, etc.
 - **Physical Layer - L1:** The lowest layer, responsible for the physical transmission of data over a medium. It includes the hardware components such as cables, switches, and network interface cards (NICs).
-

Application Layer (L5)

The Application Layer is the topmost layer in the network architecture. It provides network services directly to the user's applications. This layer is responsible for facilitating communication between software applications and lower layers of the network stack.

The Application Layer interacts with the Transport Layer (L4) to establish connections and manage data flow between applications on different devices. One can easily invent a new protocol as needed.

There are **Standard** protocols (determined by multiple organizations): HTTP, HTTPS, NFS, SSH, Bitcoin... And **Proprietary** protocols (determined by a single organization): Microsoft Exchange, Skype, Zoom, Whatsapp...

Problems In Networking - There are 2 main problems to consider when discussing networks:

- **Lossiness:** Networks are inherently lossy, meaning that data sent won't always arrive at the destination. This can be due to various reasons such as hardware failures... But more commonly, it is due to running out of *memory buffer space* in network elements.
 - **Reordering:** Networks may mess up the order of the messages sent.
-

Transport Layer (L4)

Protocols in Layer 4 provide host-to-host communication services *for processes*, they are implemented (typically) by the OS. Two most common protocols are: **TCP** (Transmission Control Protocol) and **UDP** (User Datagram Protocol).

TCP "The protocol that cares", i.e. it cares about data loss & reordering. It ensures that data is delivered reliably and in the correct order. Provides **stream** of bytes between two processes. Said to be **connection-oriented** because it establishes a connection before transmitting data. Simplified view of TCP:

- **Sender**

1. Establish a connection with the receiver by performing a *3-way handshake*.
2. Split bytes into contiguous chunks called *segments*.
3. Assign a sequence number to each segment, and keep them around until they are acknowledged by the receiver.
4. Resend segments that are not acknowledged for some time, or receiver says they are missing.
5. Slow down the sending rate if segments sent are not reaching the receiver.
6. Implement **Congestion Control** = changing the rate in response to drops.

- **Receiver**

1. Send ack messages (acknowledgments) upon receiving chunks.
2. Ask sender to resend chunks that appear to be missing for some time.
3. "Advertise" free buffer space to the sender, so it knows how much data it can send without overflowing the receiver's buffer.
4. Implement **Flow Control** (advertising) = Receiver controls sender transmission rate, thereby preventing its own buffer overflow.
5. Deliver the bytes in-order to receiving user-level app.

UDP "The protocol that doesn't care", i.e. it **does not** guarantee reliable delivery nor in-order delivery of data. Provides **datagrams** (as opposed to a stream) of bytes between two processes. Chunks might get lost or be delivered out of order, **but** per-chunk bytes integrity is supported. Said to be **connectionless** because it does not establish a connection before transmitting data. Compared to TCP, UDP is simpler and has less overhead, lower latency, and no congestion control.

Sockets Apps communicate with the Transport Layer through **socket file descriptors**, created via the `socket()` system call, instead of `open()`. A `sockfd` constitutes a communication endpoint.

Note: It is recommended to understand what IPs are before continuing this section (read next section - IP)

Ports Ports exist on the same host, and there can be a multiple of them. Each communicating process utilizes `sockfds` to identify its own communication channel, each with its own associated **port**. As IPs are not sufficient, we use a unique port for each socket, which is an **unsigned 16-bit integer** that identifies it on the host. Each chunk transmitted is associated with IP-address + port.

Ports can either be **Ephemeral** (temporary, dynamically assigned by the kernel), or **Well-Known** (static, determined by standards/known values). **Ports ≤ 1023** are reserved for well-known protocols, such as HTTP (80), HTTPS (443), etc.

Each socket is identified by a **5-tuple** consisting of:

1. Source IP address

2. Source port number
3. Destination IP address
4. Destination port number
5. Transport protocol (e.g. TCP, UDP)

Client-Server Model (TCP)

In the client-server model, a server listens for incoming connections from clients, and clients initiate connections to the server. The **server** *passively* waits for client requests and then reacts, "request-response" paradigm. The server has a well-known domain name (and/or IP address) & port. The **client** asynchronously/intermittently connects to the server, sends requests, and waits for responses. The client may have dynamic IP, and may use ephemeral port.

Sockets are initialized using values from `addrinfo` structure, which contains the following info: (1) Protocol family (e.g. IPv4, IPv6), (2) Protocol type (TCP), (3) Socket type (stream), (4) server/client.

Creating and Using TCP Sockets

server	message	client
<code>srvfd = socket(...)</code>		<code>sockfd = socket(...)</code>
<code>bind(srvfd , port)</code>		
<code>listen(srvfd)</code>		
<i>loop:</i>		
<code>clifd = accept(srvfd)</code>	← request service (transport)	<code>connect(sockfd , host+port)</code>
<code>read(clifd)</code>	← request (application)	<code>write(sockfd)</code>
<code>write(clifd)</code>	response (application) →	<code>read(sockfd)</code>

Network Layer (L3) - IP

The Network Layer is responsible for routing data between devices on different networks. It provides logical addressing and routing. The most common protocol in this layer is the Internet Protocol (IP).

Domain & Host Names Computers are associated with hierarchical human-readable "domain" names, sometimes referred to as host names, e.g. *www.google.com*. Each domain names can be backed by one or more host machines, but we will use the term host name assuming it corresponds to a single host machine.

IP Address Domain names are only for humans to read, but computers need to know the **IP address** of the host they want to communicate with. An IP address is a unique identifier assigned to each device connected to a network. Each host is mapped to a 32-bit IP address obtained via a domain name resolution protocol called **DNS** (Domain Name System). "IP" for us stands for IPv4 (Internet Protocol version 4), but there is also IPv6 which is a 128-bit address space which we will ignore.

Each IP address corresponds to a single host, unlike domain names. Domain names can point to multiple IP addresses (e.g. load balancing), while each IP address is unique to a specific host. The opposite is also not true, i.e. one host can have multiple IP addresses. For simplicity we will assume there's a 1-to-1 host-to-IP mapping.

Note: IP is used to identify host machines, while TCP/UDP are used to identify individual communication channels of processes on the host.

For example, the IP of `csa.cs.technion.ac.il` is

Notation	IP address
Decimal	2219057153 (not terribly convenient)
Binary (8x4 = 32 bits)	¹³² 10000100 ⁶⁸ 01000100 ³² 00100000 ¹ 00000001
Dot-decimal	132.68.32.1 (a bit more convenient)

Physical-Layer (L1) & Link Layer (L2)

LAN (Local Area Network) A LAN is a network that connects devices within a limited geographical area, such as a home, office, or campus. It typically uses Ethernet or Wi-Fi technologies to connect devices. Each LAN device is equipped with **Network Interface Card (NIC)**, which is a hardware component that allows the device to connect to the network.

We will focus on Ethernet, which is the most common LAN technology.

Ethernet Ethernet is a protocol that combines hardware, firmware, and software (OS). Ethernet is both a link-layer protocol (L2) as it allows nodes to communicate within the LAN by sending **frames** (byte chunks in layer L2), and a physical-layer protocol (L1) as it defines raw bits are transmitted over the physical medium.

MAC (Media Access Control) Address Each device on a LAN has a unique identifier called a **MAC address**. It is a 48-bit address that is assigned to the NIC by the manufacturer. The MAC address is used to identify devices on the same LAN and is used by Ethernet to deliver frames to the correct destination. The 48-bits are usually represented as 6 hexadecimal numbers separated by colons, e.g. `00:1A:2B:3C:4D:5E`, where the first 3 bytes are the device's manufacturer ID and the last 3 bytes are the device's unique identifier.

Ethernet Frame Structure An Ethernet frame is the basic unit of data transmission in Ethernet networks. It consists of several fields that provide information about the frame and its contents. The structure of an Ethernet frame is as follows:





Figure 5: Diagram of an Ethernet Frame Structure.


1. **Preamble:** Its purpose is purely for timing and synchronization.
2. **Start Frame Delimiter (SFD):** The actual frame starts with the very next bit.
3. **Destination/Source MAC Address:** This is the MAC of the NIC that the frame is *intended for* (Destination) and the MAC of the NIC that *sent* the frame (Source).
4. **Size:** Represents the length of the frame.
5. **data:** Contains the actual data from the upper layers, such as an IP packet from the Network Layer (L3).
6. **Cyclic Redundancy Check (CRC):** It is an **error-detection code**. i.e. the frame is dropped in case of a CRC error.
7. **Gap:** After a frame is sent, devices on the network must wait for a brief period before any device is allowed to send the next frame. This gives device interfaces time to process the frame they just received and prepare for the next one.

The Internet

IP & Routers The problem for now is LAN hardware communicate through L2, meaning two different LANs cannot communicate with each other. The solution is to use **routers**, which are devices that connect different networks and **forward** packets between them. Thus the internet is a inter-network of networks, where each inter-network is connected by routers. This is made possible using the IP (L3) protocol, which provides a way to route packets between different networks. The routers "peel off" the L2 headers of the source LAN and re-encapsulate the data within L2-headers of the destination LAN.



(a) The Internet Layers and Encapsulation.



(b) Hourglass Model of the Internet.

Hourglass Model The hourglass model is a conceptual framework that illustrates how the Internet is structured. It shows that the Internet has a narrow waist, which is the IP protocol, and that all other protocols (e.g. TCP, UDP, HTTP) are built on top of it. The model emphasizes the importance of IP as the foundation of the Internet, allowing different protocols to coexist and communicate with each other.

IP address structure An IP consists of two parts: **Subnet Part** (high order bits) and **Host Part** (low order bits). The subnet part identifies the network, while the host part identifies the specific device within that network. The subnet part is determined by the **subnet mask**, which is a 32-bit number that specifies how many bits are used for the subnet part and how many bits are used for the host part. An IP address is **always** coupled with a subnet mask; e.g. 192.168.1.1/24 means that the first 24 bits are the subnet part and the last 8 bits are the host part. **Meaning of subnet** is that all devices in the same subnet can communicate with each other directly, without going through a router. i.e. it defines the boundaries of the Ethernet LAN.

ARP (Address Resolution Protocol): $IP \rightarrow MAC$

ARP is a protocol used to map IP addresses to MAC addresses. When a device wants to communicate with another device on the same local network, it needs to know the MAC address of the destination device. If it only knows the IP address, it **broadcasts** an ARP query packet (containing the destination IP and -1 as the MAC) to the local network, asking "Who has IP address X.X.X.X? Tell me your MAC address." The device with the matching IP address responds with its MAC address, allowing the sender to encapsulate the data in a frame and send it to the correct destination. If the destination is not on the same local network, the sender will get the MAC to **its default gateway (router)**, its first-hop router, which will then forward the packet to the correct destination network.

NAT (Network Address Translation): private IPs

NAT is a technique used to allow multiple devices on a local network to share a single public IP address. NAT defines new private IP addresses for each device on the local network, allowing them to communicate with each other using these private addresses. When a device wants to communicate with the outside world, NAT translates its private IP address

to the public IP address, allowing the device to access the internet. This helps conserve the limited pool of public IP addresses and provides an additional layer of security by hiding the internal network structure from external entities. To do so, NAT maintains a mapping between pairs of (private IP addresses + port) and (public IP address + port).

DHCP (Dynamic Host Configuration Protocol): [get an IP](#)

DHCP is a **L5 protocol** used to dynamically assign IP addresses to devices on a network. The goal of DHCP server is to [allow client's host to dynamically get IP address when joining LAN](#), which can provide: (1) subnet mask, (2) Default gateway (first hop router), (3) Name + IP of DNS server. DHCP uses **UDP** to communicate allowing the ability to broadcast to all hosts in LAN. When a device connects to a network, it sends a DHCP request to the DHCP server, which responds with an available IP address and other configuration information.

Topic 7: Virtual Memory

Part II

Overall Summary

Part III

Highlights and Notes